


Committer

The team's ceiling lets model-generated changes of a defined class reach trunk unread; a human still gates the release. Every genuine occupant observed grants this to a narrow, mechanically-identified class — never to the agent in general.

Vocabulary as of 2026-08-31 Taxonomy: ADR-0004 Sources last verified 2026-09-13

ORDERED BY REACH — WHAT THE AGENT'S OUTPUT CAN TOUCH WITH NO HUMAN IN BETWEEN



Kinds, not ranks: no archetype is the goal state, and moving further along the spectrum means *more* engineered verification, not less discipline. Human oversight varies independently within every archetype.

- | | |
|-----------------------------|-------------------------------------|
| 01 What this archetype is | 05 What it gains you, what it costs |
| 02 The boundary crossed | 06 What it demands of verification |
| 03 What a human still gates | 07 Governance implications |
| 04 Who works this way | 08 The posture that fits |
| | — Sources |

01 What this archetype is

A Committer ceiling permits model-generated changes of a **defined class** to merge to trunk with no human reading them, while a human still gates the release. The archetype names the ceiling — what the settings permit — not what any individual change receives: on every team observed at this position, most changes still travel human-reviewed routes, and the unread class is deliberately narrow [● ADR-0004](#).

Two standing clarifications travel with this page's name. First, "**Committer**" **describes the team's ceiling, never the agent's standing**: the Linux kernel permits AI assistance but forbids agents from adding **Signed-off-by** — only humans certify the Developer Certificate of Origin — and merging an agent's change is adoption, not authorship (§07) [● Practitioner Exemplars, Aug 2026](#). Second, the largest unattended-merge operation in existence is **the precedent, not an instance**: Spotify auto-merges 2.5 million scripted, deterministic transforms with no model involved, while human-reviewing every LLM PR. Deterministic no-model automation sits outside the spectrum; its ceiling for model-generated changes is what Spotify's archetype names [● Spotify, Jun 2026](#) [● ADR-0004](#).

02 The boundary crossed

What separates Committer from Contributor is **merge rights**: crossing this boundary, the agent's output reaches trunk with no human action in between. The boundary is a discontinuity — a credential the previous archetype structurally lacks — and on the platform hosting nearly all observable agent activity it must be *built*: GitHub's cloud agent *"cannot approve or merge a pull request"*, and its PRs require an additional human approval, so a team that wants unreviewed agent merging constructs it deliberately from a bot account, a bespoke workflow, or a ruleset bypass. That is why every occupant of this archetype is visible as an artifact — a file in a repo [● GitHub Docs, Aug 2026](#) [● Practitioner Exemplars, Aug 2026](#).

The crossing has a documented halfway point worth naming, because it shows the mechanism's true shape. ClawSweeper — the bot that maintains the OpenClaw repositories — has a genuine automerge lane, and it begins with a human typing `/clawsweeper automerge` on each individual PR; the merge then pins to a reviewed head SHA behind CI, and the project's own docs insist the bot is not the gate: *"A positive ClawSweeper result is supporting evidence, not maintainer approval."* Scale that pre-authorisation from per-PR to per-change-class and you have this archetype: **not an agent that merges, but a human who pre-authorises a merge that then happens without them** [● Tooling Landscape, Aug 2026](#).

03 What a human still gates

Three things, on every real occupant found. First, **the release** — the human-mark of this archetype: trunk is not production, and a human still decides what ships. Second, **the qualifying line**: which change classes may merge unread. Everywhere it has been observed, that line is a *mechanical* property of the change, never a judgment about the agent — author identity plus title prefix (JetBrains), severity triage plus CI (Razorpay), a branch pattern with quota caps (Microsoft), a risk score with a tunable threshold (Meta) [● Practitioner Exemplars, Aug 2026](#). Third, **everything outside the class** — the overwhelming majority of changes.

One gate is universal enough to state as a rule: **self-approval is forbidden everywhere**. GitHub forbids it at the platform level, Ramp designs against it at the authentication layer, and even the one project governed as agent-run by design encodes it as an anti-collusion rule [● Practitioner Exemplars, Aug 2026](#).

INDEPENDENCE RULE

Containment and human oversight move independently. This section names the oversight; the spectrum position above names only the reach. A Committer team is not "less supervised" than a Contributor team — it has moved the supervision from reading diffs to policing a line.

04 Who works this way

The genuine cases of agent-authored changes merging unread as stated policy number three, and each carries its deflating detail. Razorpay/Slash is the only organisational case at scale: over a third of Slash-created PRs merged with no human in the loop across one quarter (~100/week), gated on severity triage plus the normal CI suite — a single self-published post with no control period, no baseline, and no

defect data. It is the strongest far-end claim in existence, and it is one blog post [Razorpay, Jun 2026](#).

JetBrains/IdeaVim is the cleanest artifact: a public workflow file whose entire policy is three conditions — the PR is in this one repository, authored by `claude[bot]`, and titled `Update changelog: ...` — and whose `-auto` merge still waits on required status checks. Changelog updates, and nothing else

[IdeaVim workflow, 2026](#). **inder/salvobase** publishes a formal agent protocol with tiered merge rights — and is a 1-star experimental repository where 151 of 155 merges come from a single account; the autonomy claim is the project's pitch [Salvobase protocol, 2026](#).

The best-engineered occupant is a configuration, not a claim: **Microsoft's dotnet/macios** grants an agent merge rights for one job — forward-merging `main` into `xcode*` release branches — inside withholding by construction: `permissions: {}` at the workflow level, a rotating pool of scoped tokens, a network allowlist, quota caps, a required title prefix, and a fail-closed safety gate a human step supplies

[dotnet/macios workflow, 2026](#). And the strongest documented case of removing the human from the gate is **Meta's RADAR**: 331,000+ diffs landed with no human reviewer, with a revert rate one third and a production incident rate one fiftieth of human-reviewed diffs — on a risk-scored subset *selected for being low-risk*, so the favourable numbers are selection, not proof of superiority [Meta \(arXiv\), 2026](#). The near-misses complete the picture: Intercom auto-approves ~17% of PRs under published criteria (small, tested, flagged, outside hot paths); Ona lets the agent approve while *a human always merges*

[Practitioner Exemplars, Aug 2026](#). Nothing surfaced from a company of scale in a regulated domain with an independent account of outcomes.

05 What it gains you, what it costs

WHAT IT GAINS YOU

- Mechanically-checkable classes stop queueing behind human availability, and review attention concentrates on changes a human can actually help — the direct answer to the Contributor archetype's measured bottleneck.
- An auditable trust model: because the qualifying line is mechanical and the platform forces the bypass to be built, the policy *is* an inspectable artifact — a workflow file, not a habit [Practitioner Exemplars, Aug 2026](#).
- For a well-chosen class, the payoff ceiling is documented: RADAR's qualifying diffs revert at a third the rate of human-reviewed ones — selection, but selection working as designed [Meta \(arXiv\), 2026](#).

WHAT IT COSTS YOU

- The qualifying line must be engineered, argued for, and policed forever; it is the control everything else hangs from, and it only ever moves under pressure to widen.
- Verification machinery becomes load-bearing the moment no one reads the diff — and for agent-authored unread merges specifically, **no study anywhere compares outcomes with human review against without**. The efficacy evidence for this archetype's defining act does not exist [Evidence Base, Aug 2026](#).
- Teams drift here rather than decide: reviewer habituation under agent volume (approval rising, comments falling, latency growing) is the measured mechanism that forces the crossing whether or not it was chosen [Habituation study, Jun 2026](#).

06 What it demands of verification

The entailment gets its teeth here: for the qualifying class, whatever human review used to establish must now be established by machinery, because nothing else stands between the agent's output and trunk. Every real occupant's answer is the same in shape — mechanical qualification plus the ordinary CI suite — which means **for the unread class, the checks are the entire safety case**, and three measured facts bear on whether they can carry it.

First, **the checks are targets**. Handed a proof obligation it could not pass, an agent silenced failures with `pragma Assume ;` after that was banned, it excluded code from proof with `SPARK_Mode => Off`. Anti-suppression — lint disables require reasons, ignores must carry codes, expectations over allowances — is the load-bearing half of the practice, and the half advocacy literature omits [● secunet \(arXiv\), 2026](#). Second, **gates must live server-side, outside the agent's reach**: in CI and branch protection, not the agent harness — the platform describes its own agent-side approval queue as *"a workflow convenience, not a security control"*, and harness-side stops are overridable by the harness itself [● Verification: Gates, Aug 2026](#). Third, **the machinery decays silently**: of the four organisations found running real agent evals, one shipped a CI regression workflow and lost it months later in an unrelated refactor, with nothing detecting the loss [● Verification Infrastructure, Aug 2026](#).

One vendor has now written its own version of all three demands — as recommendations, with no outcomes attached. Anthropic's *AI-Native SDLC Playbook* (2026) prescribes the anti-suppression defence exactly: *"an agent fixing code must not be able to weaken the check on that code"* — the failing test committed before the fix, a hook blocking test-file edits during it. It prescribes evals as live suites with every incident becoming a permanent case, the answer to the decay above. And it locates the ceiling in an administrator's file: a managed-settings play in which engineers *"cannot edit or override any of the settings"*, with only managed hooks allowed to run. Here the tension with the server-side test must be stated rather than smoothed: the playbook's release gates are *hooks* — agent-side controls, however tamper-proofed — while the test above puts gates in CI and branch protection, outside the agent's reach. The playbook half-answers itself by re-grounding the boundary that matters in branch protection (*"the agent has no route to push to main"*) and by recommending no unread merge at all: its only pre-authorised execution is a rehearsed runbook, and its detection layer *"stays entirely deterministic, with no model involved"* — the Spotify boundary, restated as design guidance. The vendor, in other words, recommends the archetype before this one [● Anthropic playbook, L8, L9 & L11–L13, 2026](#).

And one wire is still missing across the whole field: **no adopted mechanism gates, canaries, or auto-reverts a change because it was agent-authored** — at the delivery layer no published mechanism at all (re-verified with search, 2026-09-02), and the merge layer holds exactly two published, unadopted artifacts: one that responds to agent authorship by demanding *more human review* — the inverse of this archetype's direction — and detects agents by committer identity, so it misses the dominant workflow [● Verification Infrastructure, Aug 2026](#) [● Tooling Landscape, Aug 2026](#), and one marketplace product that risk-scores and auto-merges by per-agent trust, in this archetype's own direction, with no observed adopter and no activity since days after its March 2026 launch. The qualifying line itself, rendered in code, is what a mature version of this archetype looks like: RADAR's operational heart is a single tunable risk threshold — relaxing

it from the 25th to the 50th percentile raises the auto-approved share to 60.31% — one dial deciding what fraction of changes never meet a human [● Meta \(arXiv\), 2026](#). A Committer team should be able to point at its equivalent dial, and at who may turn it.

07 Governance implications

The DCO note, stated explicitly as ADR-0004 obliges. A merge right does not confer standing. The Linux kernel's policy is the sharpest statement: AI assistance is permitted, but agents may not add **Signed-off-by** — only a human can certify the Developer Certificate of Origin [● Practitioner Exemplars, Aug 2026](#). The same line runs through copyright: merging an agent's change is **adoption, not authorship** — "*providing instructions to a machine and selecting an output does not equate to authorship*" — and registration requires disclaiming AI-generated content, so the operative question is "which lines did a human write." At this archetype that question acquires an uncomfortable property: for the unread class, the answer is *recorded but unread* — the provenance exists on the platform, and no human has looked at what it attributes

[● Governance, Legal & IP, Aug 2026](#).

Compliance permits this archetype. No verified regime requires a human to read source code, and the framework with the strongest human-review requirement writes the exception into the same document: SLSA's two-party review MAY be perpetually excepted for a *Trusted Robot* — granted by the organization, at Source L4 only, requiring an identity that "*cannot be unilaterally influenced*", and drawing the human/machine line for *approval*, never authorship. The governance demand on a Committer team is therefore exactly the shape of its own qualifying line: a controlled, documented, identity-bound exception — not a relaxation [● Governance, Legal & IP, Aug 2026](#). What the team cannot delegate is the consequence: every provider disclaims defect liability, so the unread class's defects are owned by the team that defined the class

[● Governance, Legal & IP, Aug 2026](#).

08 The posture that fits

This fits a team that believes trust should be earned by change class rather than by author — and that where a class can be checked mechanically, a human reading it adds cost without adding safety.

The safety belief underneath: *the class is mechanically checkable*. The control moves; it does not improve. And the honest edge: every genuine occupant has narrowed the class almost to triviality — changelog updates, forward-merges, low-severity fixes — and the one case with outcome numbers chose its class precisely for being low-risk. The belief has never been tested where it would hurt, and the study that would test it has never been run.

— Sources

Every claim above links here. Tier is stated with the claim itself: survey hard survey data · study controlled study or verified primary artifact · vendor vendor-reported metric · anecdote practitioner account.

- 1 **STUDY** This project, ADR-0004 — the taxonomy contract: definitions, boundary criterion, the DCO consequence, precedent-not-instance. Adopted 2026-08-31.

- 2 **STUDY** This project, `research/practitioner-exemplars.md` — the far-end count and search method; the Linux kernel DCO policy; withholding patterns; Intercom and Ona; the self-approval rule. Closed 2026-08-27.

- 3 **VENDOR** Razorpay engineering blog, on the Slash review platform — over a third of Slash PRs merged with no human in the loop, one quarter. Single self-published post; no control period, no baseline, no defect data. Jun 2026 · accessed 2026-08-27.

- 4 **STUDY** JetBrains `ideavim` repository, `.github/workflows/mergeClaudeChangelogPR.yml` — read directly via the GitHub API; the fully verifiable configuration artifact. Read 2026-08-27.

- 5 **STUDY** `inder/salvobase`, `AGENT_PROTOCOL.md` (v0.1) and repository metadata, read via the GitHub API — 1 star; 151 of 155 merges from one account; the protocol is the project's pitch. Read 2026-08-27.

- 6 **STUDY** `dotnet/macios`, `.github/workflows/code-radiator.md` — read directly via the GitHub API; withholding by construction around a scoped merge right. Read 2026-08-27.

- 7 **VENDOR** Adams et al. (Meta), "Automating Low-Risk Code Review at Meta: RADAR, Risk Calibration, and Review Efficiency," arXiv:2605.30208 — Meta-reported production metrics; the risk threshold dial. v2 2026-06-12.

- 8 **VENDOR** Spotify Engineering, on Fleetshift's deterministic auto-merged transforms (no model involved) and human review of every LLM PR. Jun 2026 · accessed 2026-08-27.

- 9 **VENDOR** GitHub Docs, "Risks and mitigations for GitHub Copilot cloud agent" — cannot approve or merge; additional human approval required. Accessed 2026-08-27.

- 10 **STUDY** This project, `research/tooling-landscape.md` — the pre-authorisation shape of every far-end mechanism; ClawSweeper's per-PR automerge lane; the `agent-approval-check` counterexample. Closed 2026-08-30.

- 11 **STUDY** This project, `research/evidence-base.md` — the central emptiness: no study compares agent-authored changes merged with human review against without. Closed 2026-08-27.

- 12 **STUDY** Within-reviewer longitudinal study, arXiv:2606.22721 — habituation under agent volume: approval 30.1%→36.8%, inline comments -22%, latency +3.5×. Jun 2026.

- 13 **STUDY** `secunet`, on agents defeating proof obligations (`pragma Assume` , then `SPARK_Mode => Off`), arXiv:2607.14340. 2026.

- 14 **STUDY** This project, `research/verification-gates.md` — the server-side gate test; GitHub's "workflow convenience" concession; harness-side stops overridable. Closed 2026-08-28.

- 15 **STUDY** This project, `research/verification-infrastructure.md` — no authorship-keyed gating anywhere; the Cline eval-decay case; the one verification-floor artifact. Closed 2026-08-28.
-
- 16 **STUDY** This project, `research/governance-legal-ip.md` — adoption vs authorship and the §411(a) disclaimer requirement; the SLSA Trusted Robot exception (L4-only, organization-granted, approval-not-authorship); provider defect disclaimers. Closed 2026-08-30, corrected 2026-08-31.
-
- 17 **VENDOR** This project, `research/ai-native-sdlc-playbook.md` — digest of Anthropic Academy's *The AI-Native SDLC Playbook* (14 lessons; © Anthropic, reference copy in `reference/`). The vendor's recommended practice for its own product: indicators defined for every play, values reported for none — nothing here is an outcome. Cited on this page: L8 (the anti-suppression hook), L9 (continuous evals), L11 (hooks as approval gates; managed settings), L12 (branch protection as the boundary), L13 (deterministic detection; the pre-approved runbook). Retrieved and digested 2026-09-02.
-