

Glossary

The vocabulary these documents are written in. Definitions are the project's ruling usage, established against primary sources; where a common term is avoided, the reason is given.

Vocabulary as of 2026-08-31 Taxonomy: ADR-0004 Evidence in `research/vocabulary-and-disciplines.md`

01 The five archetypes

An **archetype** is a grouping of teams that work with AI in the same way — named by the team's configured *ceiling* for model-generated changes, and ordered by what the agent's output can reach with no human action in between.

Autocomplete

The agent's output reaches the open file; a human accepts every edit. The spectrum's fixed low extreme: inline AI edits with nothing else about the process changing.

Workspace

The agent holds file-write and shell tool surfaces and acts across the working tree, locally. No push credential — nothing leaves the machine without a human.

Contributor

The agent holds a push credential; its output enters the shared repository, and a human merges every change. The large middle of observed practice.

Committer

A defined class of model-generated changes holds merge rights and reaches trunk unread; a human still gates release.

Operator

Nothing human stands between the agent's change and running production — by continuous deployment from trunk or agent-held deploy credentials. Reachable per change-class; observed as no team's ceiling.

02 The spectrum's own terms

Archetype

A grouping of teams that work with AI in the same way, defined by what the human stops doing and what engineered verification replaces them with. An archetype is a team's configured *ceiling* — not what any individual change receives. Archetypes are kinds, not ranks: none is the goal state.

AVOID tier and level (imply ranking) · framework (means React or Django to this audience) · model (reads as *language model*) · stage, maturity level (imply progression)

Ceiling

The maximum delegation a team's settings and policies permit **for model-generated changes** — what an administrator flips, and what vendor terms key liability to. A property of the configured scope where the settings live (an organisation, a repository, a path); "team" is shorthand for a team at a scope. Deterministic automation with no model involved sits outside the spectrum entirely.

AVOID describing a team by its typical or modal change — that is unmeasurable from outside and drifts as volume shifts between change types

Spectrum

The set of archetypes ordered by containment ceiling — what the agent is permitted to reach. A boundary is crossed when the agent's output can reach the next surface (open file, working tree, shared repository, trunk, production) with no human action in between; the usual mechanism is a credential or tool surface, held by the agent or pre-authorised on its behalf.

AVOID maturity model, adoption curve, ladder (all imply higher is better) · write "further along the spectrum", never "higher up"

The independence rule

Containment and human oversight move independently, and the spectrum orders only containment. A team may grant an agent very little reach and supervise none of it, or grant it production reach and gate every action. Every archetype document therefore states three things separately: what the agent can reach, what a human still gates, and what verification must replace them.

AVOID writing or implying that further along the spectrum means less supervised — that conflation is the specific error these documents exist to break

Posture

The set of beliefs and constraints — about trust, risk and what AI should touch — that lands a team in a particular archetype. Each archetype page closes with the posture that fits it.

AVOID maturity, sophistication, readiness

03 **Autonomy, oversight, containment**

Autonomy

A per-action property — what this agent may do without approval, in this context. Described by naming what the human approves and when.

AVOID "L3 autonomy" or any bare level reference — no standards body has published a scale, and the two leading proposals contradict each other

Human oversight

What a human actually gates, named specifically. The register regulators use.

AVOID "human-in-the-loop" as a reassurance claim: the vendor holding the telemetry reports ~93% rising to 97% of permission prompts approved, and human review catching 13.6% of dangerous commands against 89% for automated checking — the mechanism does not imply the outcome

Containment

Restricting what an agent *can* do — sandboxes, egress controls, least-privilege tool scopes — rather than approving what it does. The field's structural answer to the limits of supervision, and the property the spectrum orders.

AVOID "guardrails" as a load-bearing noun; name the mechanism (validation, allowlist, sandbox, classifier, tripwire, permission rule, hook)

Verification

The engineered machinery that establishes a change is acceptable when a human has not read it: tests, gates, static analysis, canaries, rollback, evals.

AVOID using "review" for both human reading and automated checking — they are the two things these documents most need to keep apart

Lethal trifecta

The combination of private data access, exposure to untrusted content, and the ability to communicate externally — the configuration under which prompt injection becomes exfiltration.

Perception gap

The measured difference between how much faster developers believe AI makes them and how much faster they measurably are — consistently large, and consistently in the direction of belief exceeding measurement.

AVOID presenting it as proof that AI makes developers slower; the gap is well-evidenced, the slowdown is not

04 The disciplines

Context engineering

The discipline of assembling what an agent sees: instructions, retrieved code, tool results, history, and what is deliberately excluded. Coined by Tobi Lütke (19 Jun 2025), amplified by Andrej Karpathy (25 Jun 2025).

AVOID "prompt engineering" as the name of the discipline — it is now a subset, meaning the wording of instructions specifically

Agent harness

Everything around the model that makes it an agent: the agent loop, tool surfaces, permission modes, hooks, sandboxes, and feedback loops.

AVOID bare "harness" where eval harnesses are also discussed

Harness engineering

The discipline of designing an agent harness. Named by Mitchell Hashimoto (5 Feb 2026) and industry vocabulary within months; measurably load-bearing — harness choice moves benchmark results about as much as model choice.

AVOID "scaffolding" as a distinction to build an argument on — it is the 2024 name for the same thing

Eval harness

The evaluation infrastructure that runs a benchmark and scores results. A different thing from an agent harness; conflating them corrupts benchmark numbers.

Agent orchestration

The discipline of structuring work across multiple agent steps or agents: subagent fan-out, pipelines, gates, checkpoints, handoffs.

AVOID "graph engineering" — not a term of art anywhere in this field, and it collides with the unrelated "knowledge graph engineering"

Flow engineering

The narrower practice of encoding control flow explicitly rather than leaving the sequence to the agent's judgement. A technique within agent orchestration, not a synonym for it.

Evals

Task-specific, purpose-built measurements of whether a system does acceptable work. Distinct in practice from public benchmarks.

AVOID treating evals and benchmarks as interchangeable

Agentic engineering

Professional software development where agents do substantial work under engineered constraint — the current name for the practice the later archetypes describe.

AVOID "vibe coding" for disciplined work (it means prototypes where the diffs go unread) · "vibe engineering" (retired by its own author)

05 How these documents cite

Evidence tier

The classification every cited figure carries, stated with the claim itself — never as a blanket disclaimer: survey hard survey data · study controlled study or verified primary artifact · vendor vendor-reported metric · anecdote practitioner account.

Joint measurement

An agent benchmark score, which reflects model and harness together. Always reported as "model X + harness Y resolves N%", never as a property of the model alone.

AVOID any bare agent benchmark score; SWE-bench Verified scores entirely — deprecated by their own publisher in February 2026 for training contamination

06 A note on external scales

No standards body has published an autonomy-levels scale — verified against NIST AI RMF, ISO/IEC 22989, IEEE P3394, SAE and the EU AI Act — which is why this project defines its own archetypes and says so.

THE MILLS REVIEW IS A DIFFERENT AXIS

The FCA-commissioned **Mills Review** (Sheldon Mills, published 6 July 2026) contains an L1–L5 autonomy spectrum whose worked examples are about writing code. It orders *autonomy*; this spectrum orders *containment* — independent axes. The one rough correspondence: its L4 ("AI writes, tests and stages code; engineer approves each release") sits near the Committer archetype. Do not map the rest, and never cite its levels bare — a commissioned review is not a standard.