


Five archetypes, ordered by reach

An archetype is a team's configured ceiling for model-generated changes — what the settings and policies permit, not what any individual change receives. The spectrum orders archetypes by what the agent's output can reach with no human action in between, and each boundary is a capability the previous archetype structurally lacks. Each archetype's summary card follows; each links to its full, cited document.

Vocabulary as of 2026-08-31 Taxonomy: ADR-0004 Sources last verified 2026-09-02

ORDERED BY REACH — WHAT THE AGENT'S OUTPUT CAN TOUCH WITH NO HUMAN IN BETWEEN



Kinds, not ranks: no archetype is the goal state. Moving further along the spectrum means replacing human-in-the-loop verification with *engineered* verification — more rigor, not less. Human oversight varies independently within every archetype: every document below states separately what the agent can reach, what a human still gates, and what verification must replace them.

Autocomplete

reaches the open file

MADE OF	no tool surface — the model proposes ghost text into the open file, and a human commits every edit with a keypress. The only archetype whose ceiling is architectural rather than configured
DEMANDS FIRST	nothing new to build — the demand is the discipline of staying here deliberately, one keypress from the surfaces that act

WHAT IT GAINS YOU

- Every line is read by the person shipping it, at the moment it appears.
- The only region with clean randomised evidence in its favour — and the only one where data egress can be eliminated outright rather than managed.

WHAT IT COSTS YOU

- The smallest delta on the spectrum: the human still does everything but type.
- Reading becomes a continuous stream of micro-decisions whose acceptance rate measures satisfaction, not output.

Fits a team that believes reading every line is non-negotiable, and that anything reducing what a human reads costs more than it saves.

Workspace

reaches the working tree

MADE OF file-write and shell tool surfaces, no push credential — the agent acts across the working tree, locally; nothing leaves the machine without a human

DEMANDS FIRST containment: a sandbox, an egress allowlist, least-privilege tool scopes — which the tool field mostly does not ship by default

WHAT IT GAINS YOU

- Delegated implementation with real feedback loops — the agent edits, builds, runs and tests its own work; one human can direct several agents in parallel.
- A bounded worst case, *if the bounds are built*: local, and largely revertible.

WHAT IT COSTS YOU

- Containment becomes load-bearing exactly where the field invests least.
- The lethal trifecta — private data, untrusted content, external communication — assembles here by default unless engineered away.

Fits a team that believes risk is bounded by what the agent can reach, not by who is watching it.

[Read the full document →](#)

Contributor

reaches the shared repository

MADE OF a push credential — the agent opens pull requests under its own identity; a human merges every change. The archetype the tool market ships and the platform enforces

DEMANDS FIRST protection of the review: server-side gates, invitation-only PR creation, task-eligibility classifiers, retry caps — reviewer attention treated as an engineered budget

WHAT IT GAINS YOU

- Authoring at volume with the team's merge discipline unchanged — agent work comes through the same front door as anyone else's.
- The cheapest ceiling to hold: the boundary above is platform-enforced, so staying here means not building a bypass.

WHAT IT COSTS YOU

- The review queue becomes the bottleneck, measurably.
- Reviewer habituation is measured and worsens with the volume this archetype itself generates — the safety belief eroding under its own output.

Fits a team that believes review is the load-bearing control and should stay human — and that an agent's work should come through the same front door as anyone else's.

[Read the full document →](#)

Committer

reaches trunk, unread

MADE OF	merge rights for a defined change class, identified mechanically — author identity, title prefix, severity, risk score — which reaches trunk with no human reading it; a human still gates release
DEMANDS FIRST	the qualifying line rendered in code, anti-suppression on the checks, gates enforced server-side, and revert machinery for when the line is wrong

WHAT IT GAINS YOU

- Mechanically-checkable classes stop queueing behind human availability; review attention concentrates where a human can actually help.
- An auditable trust model — every occupant's policy is an inspectable artifact, a workflow file rather than a habit.

WHAT IT COSTS YOU

- The qualifying line must be engineered, argued for, and policed forever.
- Verification is load-bearing the moment no one reads the diff — and no study anywhere tests unread agent merges against reviewed ones.

Fits a team that believes trust should be earned by change class rather than by author — and that where a class can be checked mechanically, a human reading it adds cost without adding safety.

[Read the full document →](#)

Operator

reaches running production

MADE OF	a pre-authorised path to production — nothing human stands between the agent's change and running prod, whether continuous deployment carries an unread merge or the agent holds deploy credentials
DEMANDS FIRST	the full engineered chain — gates, canaries, auto-revert, provenance — trusted <i>before</i> the path opens; almost none of it exists wired to autonomy anywhere

WHAT IT GAINS YOU

- Full pipeline speed for the change classes that qualify.
- The clearest possible safety case: stated in machinery, not in habit.

WHAT IT COSTS YOU

- First-of-kind wiring, built and policed by you; liability keys to the setting.
- Observed as no team's ceiling — reached per change-class, by narrowing what qualifies, never by widening who may act.

Fits a team that believes engineered verification can be made a stronger check than a human in the path — and that the honest route there is to narrow what qualifies, never to widen who may act.

[Read the full document →](#)

Every claim in the five documents carries a citation chip with its source, date, and evidence tier — hard survey data, controlled study or verified primary artifact, vendor-reported metric, or practitioner account — stated with the claim itself. The glossary holds the ruling vocabulary these pages are written in.